

Approach

1. Data Understanding

a. Data Structure

Train data set contains 20800 rows, and 5 columns

Id: unique identifier

Title: title of the article

Author: author of the article

Text: article

Label: 1 for fake news, 0 for real news

b. Missing Values

- i. There were missing values in both test and train dataset in text columns

```
In [44]: ## Check for missing values
train.isnull().sum()
```

```
Out[44]: id          0
         title      558
         author    1957
         text       39
         label      0
         dtype: int64
```

Author column has more than 5% missing values - cannot drop all rows. Imputing with empty text

```
In [45]: test.isnull().sum()
```

```
Out[45]: id          0
         title     122
         author     503
         text        7
         dtype: int64
```

Imputed empty string “ ” to treat missing values.

c. Data Cleaning:

- i. Removing punctuations and digits and stopwords (from nltk) from text
- ii. Word tokenization and Lemmatizing nouns, to capture action words for feature engineering

d. Feature Engineering

- i. Quantifying the following hypothesis about fake news
 1. Polarizing words
 2. Vague and opinionated - less citations
 3. Shorter due to lack of citations and supporting facts
- ii. Creating below variables
 1. Title_polarity : polarity sentiment of title
 2. Text_polarity: polarity sentiment of text
 3. Ttl_words: count of total meaningful words in text
 4. Text_digit_cnt: count of numerical words in text eg: “1984”, “67” etc

2. Exploratory Data Analysis

a. Text Columns

- i. Word Cloud words in text, author and title column for test and train
- ii. Word Cloud of words in text, author and title column for fake and real news in train
 1. Fake news articles have more anonymous and vague named bloggers as authors and use more trending names

Fake news authors:



Real new authors:



b. Numeric Data Columns

- i. Title polarity was dropped, as it did not capture anything
- ii. The distribution and features of these values align with the hypothesis

```
In [96]: ## Analyzing new variables distribution w.r.t. to the label variable
print(pd.pivot_table(train, index = 'label', values = cols_num[1:],aggfunc = 'sum'))
print(pd.pivot_table(train, index = 'label', values = cols_num[1:],aggfunc = 'min'))
print(pd.pivot_table(train, index = 'label', values = cols_num[1:],aggfunc = 'max'))
```

	text_digit_cnt	text_polarity	ttr_wrds
label			
0	108444	631.065967	9522464
1	79922	589.091514	6807367

	text_digit_cnt	text_polarity	ttr_wrds
label			
0	0	-0.45	1
1	0	-1.00	0

	text_digit_cnt	text_polarity	ttr_wrds
label			
0	235	0.65	15349
1	470	1.00	24870


	text_digit_cnt	text_polarity	ttr_wrds
label			
0	10.440358	0.060755	916.767498
1	7.675214	0.056573	653.737348

On average fake news articles have less data citations, range to extreme polarity, and are shorter than real news

Visualization of text sentiments with count of words and label

```
In [57]: ## text polarity and ttr_wrds have small positive correlation, check their distribution
plt.scatter(train.text_polarity,train.ttr_wrds ,c = train.label,label = 'label',alpha = 0.5)
plt.legend()
```

Out[57]: <matplotlib.colorbar.Colorbar at 0x201d3299c40>



3. Modelling

a. Data Preparation

- i. Text Columns
 1. Used TF-IDF vectorizer to extract features from text columns
 2. For title and text ignored the words appearing in less 5% of data
 3. For author, vectorized all data
- ii. Numerical Columns
 1. Min Max Scaler to scale the data
- iii. Fitted the tfidf vectorizer on entire (train + test) dataset

b. Model Selection

- i. Binary Classification problem, picked 4 models:
 1. Logistic Regression
 2. Random forest Classification

3. Stochastic Gradient Descent Classifier
4. Support Vector Classifier

c. Model Testing

- i. 5 fold cross validation score
- ii. Train and test split : 75:25
 1. Metrics: Precision, Recall, F1 Score
 2. Analysed ROC Curve

Model	Cross Validation	Metrics on 25% test
Stochastic Gradient Descent Classifier	[0.9875, 0.98846154 , 0.9900641, 0.98173077, 0.98910256]	Confusion Matrix: [[2605 37] [30 2528]] Precision score: 0.9856 Recall score: 0.9883 F1 score: 0.9869
Logistic Regression	[0.98044872, 0.97980769, 0.98397436, 0.975, 0.98044872]	Confusion Matrix: [[2598 44] [49 2509]] Precision score: 0.9828 Recall score: 0.9808 F1 score: 0.9818
Random Forest Classifier	[0.94807692, 0.94871795, 0.94230769, 0.94775641, 0.94391026]	Confusion Matrix: [[2528 114] [164 2394]] Precision score: 0.9545 Recall score: 0.9359 F1 score: 0.9451
SVC	[0.98878205, 0.99166667, 0.99134615, 0.98589744, 0.99038462]	Confusion Matrix: [[2614 28] [24 2534]] Precision score: 0.9891 Recall score: 0.9906 F1 score: 0.9898

d. Conclusion

- i. ROC curve: SVC and SGD perform better
- ii. F1 Score: SVC and SGD perform better

e. Final Predictions (tested on Kaggle, SVC and SGD performed better than other two models)

Submission and Description	Private Score	Public Score	Use for Final Score
SVC(random_state42).csv 3 minutes ago by Apurva S svc clf	0.99203	0.98589	<input type="checkbox"/>
SGDClassifier(random_state42).csv 3 minutes ago by Apurva S sgdc clf	0.98956	0.98141	<input type="checkbox"/>
RandomForestClassifier(random_state42).csv 4 minutes ago by Apurva S random forest clf	0.95714	0.95000	<input type="checkbox"/>
LogisticRegression(random_state42).csv 4 minutes ago by Apurva S Log CLf	0.98543	0.97628	<input type="checkbox"/>